



Y.A.S.K.

Yet Another Stencil Kernel

Chuck Yount, Principal Engineer
Intel Corporation

April 17, 2017

chuck.yount@intel.com

Outline

Introduction

- Overview
- Example stencils and performance on Intel® Xeon® and Intel® Xeon Phi™ processors

Using YASK

- Build and test
- Output
- Use model
- Run-time options: hierarchy sizes, MPI, wave-front blocking
- Stencil, vectorization, loop, and advanced customization
- Collaboration opportunities

YASK Software architecture

- Vector folding and the fold builder
- Loop-code generator
- Memory accessor
- Debug output

Introduction to YASK

Overview

YASK: Yet Another Stencil Kernel

- Goal: facilitate exploration of the stencil-performance design space for Intel® Xeon® or Intel® Xeon Phi™ processors supporting the AVX, AVX2, or AVX-512 instruction sets

Features

- Supports trade-off studies for coding options for
 - Vector-folding
 - Cache blocking
 - Memory layout
 - Loop construction
 - Temporal wave-front blocking
 - MPI halo exchanges
 - And more
- Is a collection of C++ code, code-generators and other scripts

Example 1: Iso3dfd stencil



Description

- Finite-difference code found in seismic-imaging software used by energy-exploration companies to predict the location of oil and gas deposits

Performance* on 1024^3 problem size per node

- Intel® Xeon Phi™ processor 7250 ("Knights Landing")
 - 17.5 GPoints/sec on one node (68 cores)
- Intel® Xeon® processor E5-2697 v4 ("Broadwell")
 - 5.58 GPoints/sec on one node (2 sockets of 18 cores each)
- Intel® Xeon® processor E5-2699 v3 ("Haswell")
 - 4.71 GPoints/sec on one node (2 sockets of 18 cores each)

*Observed performance for illustration and comparison; not guaranteed.

Image from <https://commons.wikimedia.org/wiki/File:PlatformHolly.jpg>. Public domain--U.S. DoE.

Iso3dfd build and run “recipes”

Intel(R) Xeon Phi(TM) CPU 7250 @ 1.40GHz, 16GB MCDRAM flat mode

- `make clean; make stencil=iso3dfd arch=knl`
- `./stencil-run.sh -arch knl -- -d 1024`
...
`best-throughput (prob-size-points/sec): 17.5395G`
`best-throughput (point-updates/sec): 17.5395G`
`best-throughput (est-FLOPS): 1.06991T`

2-socket Intel(R) Xeon(R) CPU E5-2697 v4 @ 2.30GHz

- `make clean; make stencil=iso3dfd arch=hsw`
- `./stencil-run.sh -arch hsw -ranks 2 -- -d 1024 -dx 512`
...
`best-throughput (prob-size-points/sec): 5.58363G`
`best-throughput (point-updates/sec): 5.58363G`
`best-throughput (est-FLOPS): 340.602G`
- Note that half of the x-domain-size is specified for each of the two ranks so that the overall size is 1024^3
- The same recipe may be used for the Intel® Xeon® CPU E5-2699 v3 @ 2.30GHz

Example 2: AWP stencil

Description


- Primary compute kernel for the Anelastic Wave Propagation—ODC earthquake simulator: <http://hpgeoc.sdsc.edu/AWPODC>
 - Consists of 26 grids in a staggered-grid formulation

Performance* on $1024^2 \times 128$ problem size per node

- Intel® Xeon Phi™ processor 7250 (“Knights Landing”)
 - 23.1 GPoints/sec on one node (68 cores)
- Intel® Xeon® processor E5-2697 v4 (“Broadwell”)
 - 5.58 GPoints/sec on one node (2 sockets of 18 cores each)
- Intel® Xeon® processor E5-2699 v3 (“Haswell”)
 - 4.82 GPoints/sec on one node (2 sockets of 18 cores each)

*Observed performance for illustration and comparison; not guaranteed.

Image from https://commons.wikimedia.org/wiki/File:Izmit_11-12-99.gif. Public domain--U.S.G.S.



AWP build and run “recipes”

Intel® Xeon Phi™ CPU 7250 @ 1.40GHz, 16GB MCDRAM flat mode

- `make clean; make stencil=awp arch=knl`
- `./stencil-run.sh -arch knl -- \`
`-dx 1024 -dy 1024 -dz 128`
...
`best-throughput (prob-size-points/sec): 1.54397G`
`best-throughput (point-updates/sec): 23.1596G`
`best-throughput (est-FLOPS): 416.872G`

2-socket Intel® Xeon® CPU E5-2697 v4 @ 2.30GHz

- `make clean; make stencil=awp arch=hsw`
- `./stencil-run.sh -arch hsw -ranks 2 -- \`
`-dx 512 -dy 1024 -dz 128`
...
`best-throughput (prob-size-points/sec): 372.652M`
`best-throughput (point-updates/sec): 5.58977G`
`best-throughput (est-FLOPS): 100.616G`
- Note that half of the x-domain-size is specified for each of the two ranks so that the overall size is $1024^2 \times 128$
- The same recipe may be used for the Intel® Xeon® CPU E5-2699 v3 @ 2.30GHz

Using YASK

Initial build and test

Install

- Download the code from the 'GIT REPO' link at <https://01.org/yask>
- Install all the prerequisites from the README file

Build and run the default test program

- Type 'make -arch *name* -stencil *name*' per the README file
- Run the program using the stencil-run.sh script
 - Run natively on Xeon™ or Xeon Phi™ processors (KNL)
 - Use the -mic option to run on a Xeon Phi™ coprocessor (KNC)
 - Run under SDE to emulate hardware you don't have
- If it doesn't build and/or run, check the prerequisites

Typical run and output

Settings are printed

- Stencil name
- Hierarchical sizes (spatial and temporal): rank, region, block, cluster, and vector
- Other miscellaneous compile-time and run-time settings

A number of trials (default=3) is run

- Time and throughput (points per sec) are printed
- FP-rate is estimated based on number of FP ops in original scalar spec
- Best result across the trials is re-printed

Validation

- If the '-v' option was used, a non-vectorized, non-tiled version of the code is run, the results are compared, and 'PASSED' or 'FAILED' is printed
- Validation is slow; run with a small problem size!
- If you get near-miss errors during validation, it may be due to rounding error instead of a bug; try building with "real_bytes=8" to check

Use model


Goal

- YASK is a tool for exploring the stencil design space
- It is not a library (yet)

Typical usage model

- Identify stencil(s) used in your application
- Use existing stencil(s) in YASK or write your own
- Use YASK to find well-performing parameters
- Integrate the stencil code back into your application

Tiling applied at multiple levels



Common run-time options

Settings controlled from the 'stencil-run.sh' script

- Binary selection via 'arch' option
- Number of MPI ranks
- Which Xeon Phi coprocessor or other host to use
- Run with '-h' option to get help

Settings controlled from the 'stencil.arch.exe' binary (passed through from the stencil-run.sh script)

- Spatial/temporal rank-domain size (overall problem if not using MPI): -d*
- Spatial/temporal region size (used for temporal wave-front blocking): -r*
- Block size: -b*
- Sub-block size: -sb*
- Padding: -p*
 - Used to fine-tune data alignment across rows and columns
 - The specified value is added to the halo size during memory allocation
- Number of trials: -t
- Enable validation: -v
- Run with '-h' option to get more help on these and other options

Enabling MPI

Scope

- The MPI implementation in YASK exchanges halos with all neighbors in all directions
 - For 3D problems, this can be up to 26 neighbors (3^3-1)
 - For 4D problems, this can be up to 80 neighbors (3^4-1)

Usage

- Compile with MPI enabled using “make mpi=1 ...”
- Prefix “stencil-run.sh” with the appropriate MPI command, e.g., mpirun -n 4 -ppn 1
 - Use the “-ranks <n>” option to “stencil-run.sh” as a shortcut to run more than one rank on a single node
- Note: the -d* options control the rank size, so the overall problem size increases by the number of ranks (weak scaling)

Enabling temporal wave-front blocking

Purpose

- The temporal blocking in YASK is designed to exploit large shared caches, e.g., when using the Intel® Xeon Phi™ processor in MCDRAM cache mode

Usage

- Temporal wave-front blocking is done using the “regions” level of the hierarchy shown earlier
 - Spatial blocks within each region are evaluated in parallel using OpenMP
 - The time slices within each region are evaluated sequentially to reuse memory
 - Regions are evaluated sequentially to increase shared-cache locality
- Executable run-time options
 - `-rt <n>` sets the number of time slices in each region
 - `-r* <n>` sets the spatial size of each region
 - Example: `stencil-run.sh -arch knl -d 1920 -dt 50 -r 768 -rt 25`
 - Note: the default setting of `-rt` is one (1), and the default setting of `-r` is zero (0), which means the region size is the same as the rank size.

Stencil customization

Stencil Type

- Use the `'stencil=stencil-name'` argument to the `make` command to select a stencil (required)
 - The *stencil-name* string is passed to the `foldBuilder` tool
- Examples of current provided stencils
 - `'iso3dfd'`: an isotropic acoustic wave equation
 - `'3axis'`, `'9axis'`, `'3plane'`, and `'cube'`: common 3D symmetric shapes (as used in the [vector-folding paper](#))
 - `'ave'`: the simple 27-pt stencil from the miniGhost benchmark
 - `'awp'`: a simplified version of [AWP-ODC](#) earthquake simulation stencils
- Write your own by modifying code in `src/foldBuilder`
 - Implement the `StencilBase` interface using the `stencils/*Stencil.hpp` files as examples

Stencil customization (cont.)

Stencil size

- Use the `'radius= n '` argument to the make command
 - The n value is passed to the `'foldBuilder'` tool
 - Default is different for various stencils and ignored for some
 - See Makefile for current defaults
- Write your own by modifying code in `src/foldBuilder`
 - Follow the existing examples to pass the `'radius'` parameter to your stencil code if applicable

Other parameters

- If you're developing your own stencil, you can add more parameters similar to the `'radius'` one

Stencil customization (cont.)

Advanced

- The provided stencils assume uniformity across the entire 3D grid
 - The 'foldBuilder' tool evaluates the stencil code only from the origin to the extent of a vector
- Some stencil applications require special code at boundaries or other conditions
 - A mechanism for handling conditions that specify sub-domains is under development
 - See the awp_elastic stencil for an example
 - Since this is under development, the syntax may change

Vector-folding introduction

Concept

- Store small 2D or 3D block of data into each vector
- Pros: reduces memory BW requirements compared to traditional 1D in-line vectors
- Cons: requires data pre-conditioning (element rearranging) and additional shift and blend operations preceding SIMD arithmetic operations

Results

- Significant speedup shown on Intel® Xeon Phi™ Coprocessor
- Combining with loop tiling enables even more speedup

For more information


- Refer to paper on [Vector Folding from HPCC 2015](#)

Traditional in-line 1D vectorization

25-point 3D stencil
input vectorized using
8-element vectors,
each parallel to x-axis



Inner 3D loop iterates
in x direction, i.e.,
same dimension as
vectorization



Previous
iteration

Current
iteration

Need to read 17 new cache
lines* for each iteration (8 of
the 25 vectors overlap in x
dimension)




One 8-element
vector output per
iteration

Steady-state memory BW =
17 new cache lines input to
calculate each vector of output

*Assuming cache line size = vector size.

Reduce BW via vector folding

25-point 3D stencil
input vectorized using
8-element vectors,
*each containing a 4x2
grid in the x-y plane*



Need to read only 7 new
cache lines for each iteration
(vectors overlap in x-y
dimensions within an iteration
and in z dimension between
iterations)



Inner 3D loop iterates in z
direction, i.e., *perpendicular*
to 2D vector



One 8-element (4x2)
vector output per
iteration

Steady-state memory BW = **7**
new cache lines input to
calculate each vector of output:
2.4x lower than in-line

Vector-folding customization

Vector fold

- Use the `fold='x=n,y=n,z=n'` argument to the `make` command to control how much vectorization is done in each dimension
 - The values are passed to the 'foldBuilder' tool
 - Example: `make fold='x=1,y=2,z=8'` generates code using a 1x2x8 fold
 - Make sure the product of the fold lengths equals the vector size of the target architecture and FP precision (single or double)

Vector cluster

- Use the `cluster='x=n,y=n,z=n'` argument to the `make` command to control how many vectors are calculated and output simultaneously
 - The values are passed to the 'foldBuilder' tool
 - The default is 1x1x1, or one HW vector
 - This essentially implements loop unrolling in multiple dimensions

Loop-structure customization

The 'gen-loops.pl' script creates the loop-control code

- There are 4 loop-control codes
 - 'Outer' loops break the whole problem into OpenMP regions (typically, only one OpenMP region is used)
 - 'Region' loops break each OpenMP region into cache blocks
 - 'Block' loops iterate over each vector cluster in a cache block
 - 'Halo' loops are used for copying data for halo exchanges

Usage

- See the Makefile for default invocations or run 'make -n'
- Run './gen-loops.pl' without any parameters to get help on more options: index ordering, OpenMP scheduling, etc.
- Run the script before the make command or specify the *LOOP_ARGS variables in the make command to override

Misc. advanced customization

More compile-time options to the make command

- Use `'real_bytes=n'` to set the size of a float: *n*=4 for single-precision or *n*=8 for double-precision (default=4; 8 for 'ave' stencil)
- Use `'EXTRA_MACROS='macro-settings'` set other CPP macros
 - `'PFDL1=n1 PFDL2=n2'` to change the prefetch distances; only used in the prefetch code generated from 'gen-loops.pl', not in compiler-generated prefetch code
 - Example: `'make EXTRA_MACROS='PFDL2=15''`
 - See *.hpp for macro definitions
- Run `'make echo-settings'` to see other make variables

Automatic Tuning

YASK includes a genetic-algorithm-based automatic tuning utility

- Tunes compile-time and run-time options and parameters to maximize throughput
- Run `'stencil-tuner.pl -h'` to get tuner options
- Example run:
 - Tune the parameters for the `'iso3dfd'` stencil on the KNL architecture, constraining the memory usage between 14 and 16 GiB
 - `./stencil-tuner.pl -stencil=iso3dfd -arch=knl -mem=14-16`
- Results
 - The settings and throughput of each trial is saved into a `.csv` file
 - Use the `'stencil-tuner-summary.csh'` script to find the best result
 - Can use after the tuner exits or anytime while it's running
 - Each result includes the `"make"` and `"run"` commands
 - Often these commands will include parameter settings which are not strictly necessary to obtain the best results
 - You may wish to experiment manually with the best outcome to find simpler commands that produce similar results

Collaboration

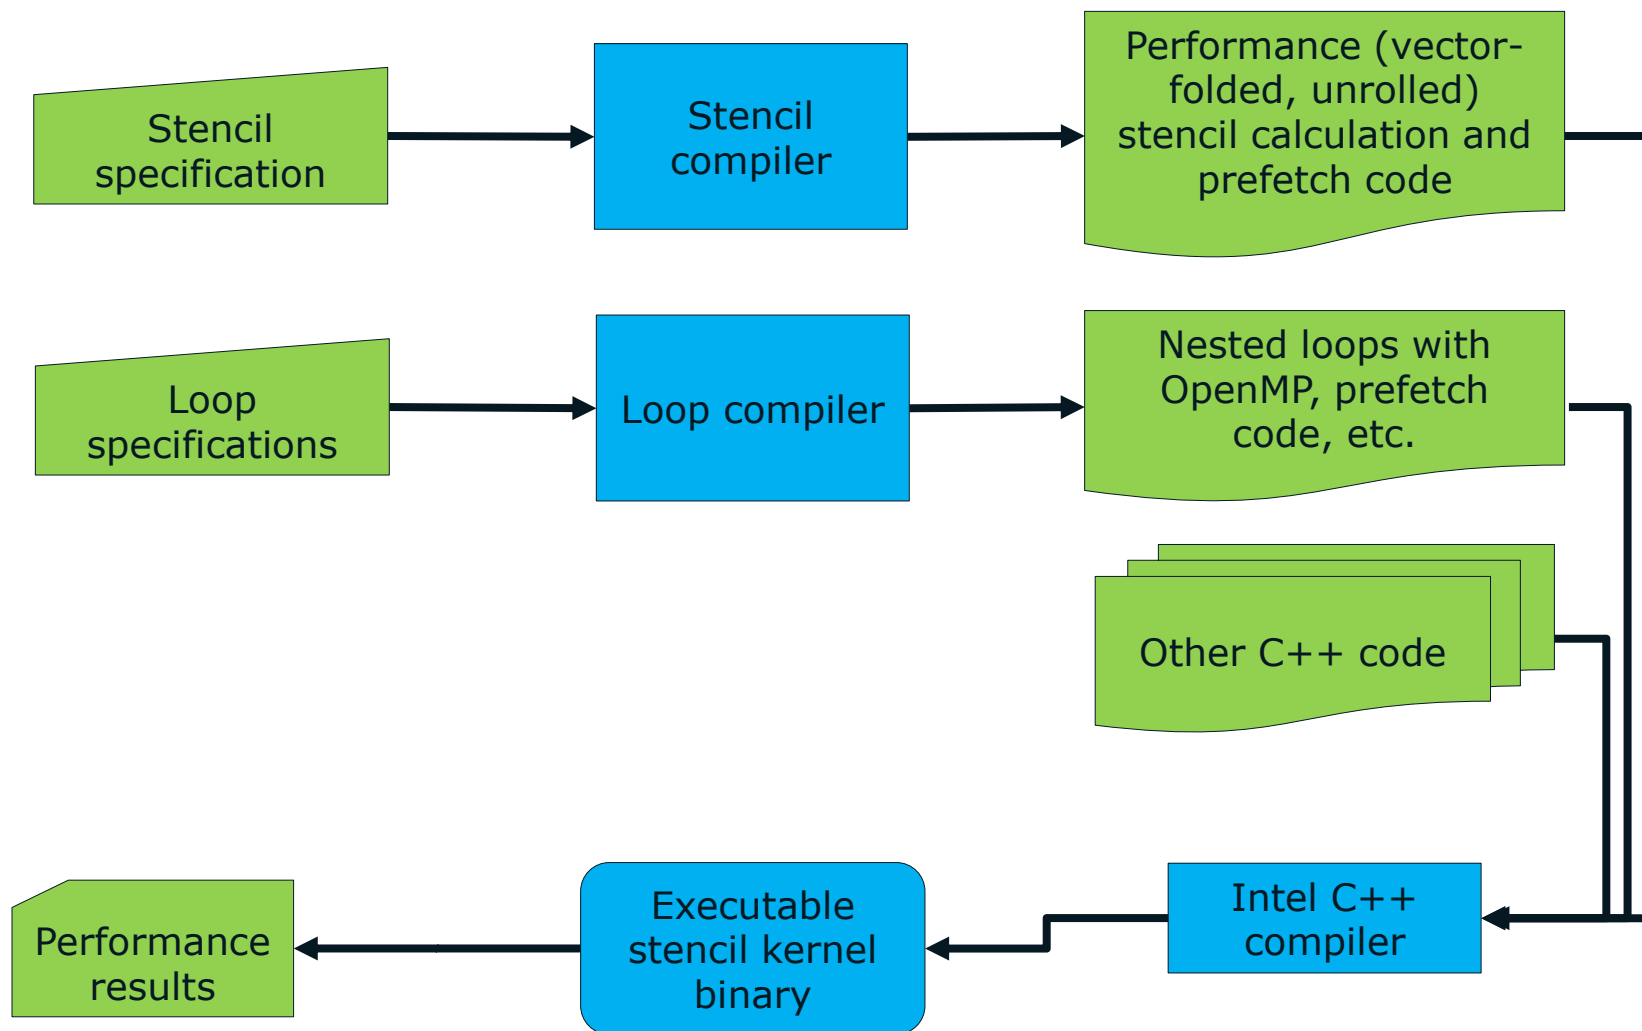
Use the blog at <https://01.org/yask> to ask and/or answer questions

Submit useful changes for review via github

Contact the author of this presentation for further collaboration opportunities

YASK Software Architecture

High-level components



Fold-builder code generator

Goal: automate the tedious and error-prone process of creating high-performance stencil code

Input

- Inherit from a C++ abstract 'StencilBase' class to create a new stencil type
- Define the grid(s) to be used and the names of their dimensions, e.g., "t", "x", "y", "z"
- Implement the 'define' method to define how one point in each grid is calculated from others
- Use loops, functions for coefficients, recursion, etc.

Process

- Compile code into fold-builder executable
- Run executable, specifying any stencil parameters (e.g., order), target architecture, etc.
- Code generator evaluates the 'define' method to create an abstract syntax tree (AST)
- AST is traversed, and optimized code is output

Output

- Efficient function to calculate stencil
 - Unrolled loops, intrinsics to construct unaligned vectors of points, etc.
 - Calls to memory accessor object
- Functions for prefetching to L1 and L2

Example input stencil code

```
class Iso3dfdStencil : public StencilOrderBase {
...
  INIT_GRID_4D(pressure, t, x, y, z),
  INIT_GRID_3D(vel, x, y, z)
...
  virtual void define(const IntTuple& offsets) {
...
    // start with center value multiplied by coeff 0.
    GridValue v = pressure(t, x, y, z) * coeff(0);

    // add values from x, y, and z axes multiplied by the
    // coeff for the given radius.
    for (int r = 1; r <= _order/2; r++) {

      // Add values from axes at radius r.
      v += (
        // x-axis.
        pressure(t, x-r, y, z) +
        pressure(t, x+r, y, z) +

        // y-axis.
        pressure(t, x, y-r, z) +
        pressure(t, x, y+r, z) +

        // z-axis.
        pressure(t, x, y, z-r) +
        pressure(t, x, y, z+r)

      ) * coeff(r);

    }

    // finish equation, including t-1 and velocity components.
    v = (2.0 * pressure(t, x, y, z))
      - pressure(t-1, x, y, z) // subtract pressure from t-1.
      + (v * vel(x, y, z));    // add v * velocity.

    // define the value at t+1 to be equal to (really, an approximation of) v.
    pressure(t+1, x, y, z) EQUALS v;
  }
}
```

Declare 2 grids: 4D
"pressure" and 3D
"vel"

Write function to
define equation for
"pressure" update

The final equation
uses declarative (not
imperative) style via
the EQUALS macro

Loop-code generator

Script that generates code for nested loops

- Input: Very simple DSL (domain-specific language)
 - `omp loop(y) { loop(x, z) { calc(stencil); } }`
 - Can easily change loop types, index ordering
- Output: C++ code to be included in function bodies
 - Loops annotated with OMP as requested
 - Inner loop might generate several loops, e.g.,
 - Prefetch L2
 - Prefetch L1
 - Calculate stencil and prefetch L2 and L1
 - Calculate and prefetch L1 only to avoid over-prefetching L2

Memory accessor

- C++ classes to allocate and access 3D & 4D matrices of vectors of floats or doubles
 - Construction specifies 'w, x, y, z' dimensions and padding sizes; padding includes halos (use 'w' for 4D problems)
 - Read and write access via methods: per vector for highest speed; per element for scalar code.
- Actual memory layout is encapsulated and defined via inheritance
 - Map indices to 1-D mem address
 - 24 simple permutations of minor-to-major ordering in 4D
 - More complex mappings possible, e.g., tiling, space-filling curves
 - 't' dimension is used for time and is template for performance
- Uses concrete inheritance to allow inlining
 - Gives compiler full access to memory-layout formula
 - Allows common sub-expression elimination and other optimizations

Debug features

Can enable or disable various output by setting macros and rebuilding, e.g.,

- TRACE: print each stencil calculation
- TRACE_MEM: print each matrix read, write, prefetch, eviction
- TRACE_INTRINSICS: print before-and-after each permutation

Built-in memory-access tracker

- Models an infinite L1 or L2 cache
- Tracks reads, writes, prefetches, evictions
- Reports any un-prefetched read or un-read prefetch
- Reports summary stats
- Very useful for debugging prefetch code

Example cache-model output

```
modeling cache...
cache L2: redundant prefetch of 0x2aaabfa45a40 at line 193 after a read at line 85.
cache L2: redundant prefetch of 0x2aaabfa45a80 at line 193 after a read at line 85.
cache L2: redundant prefetch of 0x2aaabfa45a40 at line 195 after a prefetch at line 193.
cache L2: redundant prefetch of 0x2aaabfa45a40 at line 196 after a prefetch at line 195.
...
done modeling cache...
cache L2: read of 0x2aaabf9c3240 from line 85 without any eviction.
cache L2: read of 0x2aaabf9c3280 from line 85 without any eviction.
...
cache L2: prefetch of 0x2aaabfa53b00 from line 318 without any read.
cache L2: prefetch of 0x2aaabfa53b40 from line 318 without any read.
...
cache L2 stats:
  cur size = 324714 lines (19.8190 MB).
  max size = 324714 lines (19.8190 MB).
  ave size = 185126 lines (11.2992 MB).
  num reads = 722400.
    num reads of missing lines = 0.
    num lines read but never evicted = 321700.
  num prefetches = 1458800.
    num prefetches of lines never subsequently read = 3014.
    num prefetches of lines already in cache = 404686.
  num evictions = 0.
    num evictions to non-existent lines = 0.
  num prefetches into L1 = 729400.
    num prefetches into L1 of missing lines = 0.
```

Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

